

Automatic Differentiation as a Tool for Computational Science

Laurent Hascoët
Paul Hovland
Jan Hückelheim
Sri Hari Krishna Narayanan

SIAM AN24 Student Days



U.S. DEPARTMENT OF
ENERGY

Office of
Science

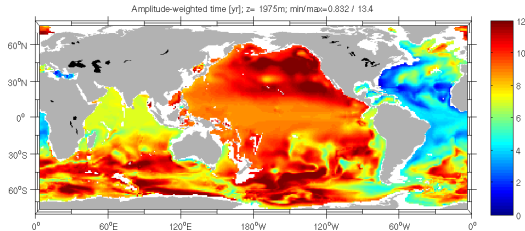
What do we have in store

- ▶ Presented Today:
 - ▶ Introduction
 - ▶ Demo & Hands on: AD basics
 - ▶ Demo & Hands on: Using AD for optimization
- ▶ Further Material:
 - ▶ Seed matrices
 - ▶ Memory requirements
 - ▶ AD for parallel programs
 - ▶ Know what you are differentiating
 - ▶ Adding AD to existing code

Resources: <https://tinyurl.com/siaman24ad>

Why derivatives?

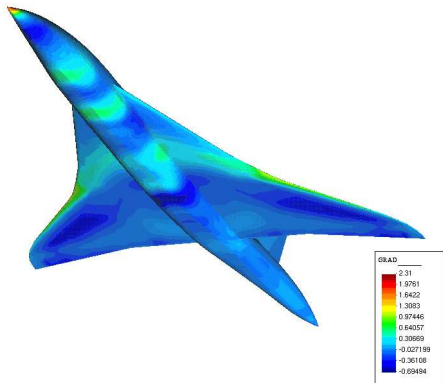
Sensitivity Analysis:



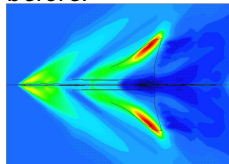
Find sensitivity of the computed field *wrt* one input parameter

Why derivatives?

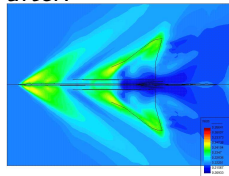
Gradient-driven optimization



before:




after:



Why derivatives?

Inverse problems:

from measurements and model, estimate hidden parameters



Other uses: Reduced models, Error estimation, Mesh adaption, Uncertainty Quantification, Backpropagation for ML training. . .

What is AD?

The chain rule applied to algorithms

See any algorithm/program $P: \{I_1; I_2; \dots I_p; \}$ as:

$$F: \mathbb{R}^n \rightarrow \mathbb{R}^m \quad F = f_p \circ f_{p-1} \circ \dots \circ f_1$$

Define for short:

$$V_0 = \mathbf{input} \quad \text{and} \quad V_k = f_k(V_{k-1})$$

Apply the chain rule:

$$F'(V_0) = f'_p(V_{p-1}) \times f'_{p-1}(V_{p-2}) \times \dots \times f'_1(V_0)$$

...and transform P to make it compute that.

Cost considerations

$$F'(V_0) = f'_p(V_{p-1}) \times f'_{p-1}(V_{p-2}) \times \cdots \times f'_1(V_0)$$

is often expensive:

- ▶ in computation time
- ▶ in storage space

What can save us:

- ▶ The shapes of the f'_k matter
- ▶ The final usage may not require the full F' but only a projection

Classical projections of F'

- ▶ $F' \times \dot{V}_0$, “forward” or “tangent” mode
- ▶ $\overline{V}_p \times F'$ “reverse” or “adjoint” mode

When full F' needed, use multi-directional AD
→ $F' \times Id$ (or $Id \times F'$),
→ possibly compressed as $F' \times S$ (or $S \times F'$)

For higher-order derivatives, differentiate F'
If directional, differentiate $F' \times \dot{V}_0$

Demo (with Tapenade)

rosenbrock.f90 :

```
REAL*8 FUNCTION ROSENBROCK(x,n) RESULT(y)
  INTEGER :: n
  REAL*8 :: x(0:n)
  y = SUM(100.d0*(x(1:n) - x(0:n-1)**2)**2
          + (1-x(0:n-1))**2)
END FUNCTION ROSENBROCK
```

- ▶ $n + 1$ inputs x , scalar output y
- ▶ Looking for $\frac{dy}{dx}$

Demo with Tapenade: tangent

```
$ tapenade rosenbrock.f90  
$> -head "rosenbrock(y)/(x)" -d
```

rosenbrock_d.f90 :

```
arg1d(:) = 100.d0*2*(x(1:n)-x(0:n-1)**2)  
          *(xd(1:n)-2*x(0:n-1)*xd(0:n-1))  
          - 2*(1-x(0:n-1))*xd(0:n-1)  
arg1 = 100.d0*(x(1:n) - x(0:n-1)**2)**2  
       + (1-x(0:n-1))**2  
yd = SUM(arg1d(:))  
y = SUM(arg1(:))
```

driverTgt.f90 initializes xd to 0,0,0,0,1,0,0,0,0,0 (n=9)

```
$ gfortran rosenbrock_d.f90 driverTgt.f90 -o tangent.exe  
$ ./tangent.exe  
Rosenbrock tangent: -260.4622
```

Demo with Tapenade: gradient

```
$ tapenade rosenbrock.f90  
$> -head "rosenbrock(y)/(x)" -b
```

rosenbrock_b.f90 :

```
xb = 0.0_8  
tempb = 2*(x(1:n)-x(0:n-1)**2)*100.d0*yb  
xb(0:n-1) = xb(0:n-1) - 2*(1-x(0:n-1))*yb  
xb(1:n) = xb(1:n) + tempb  
xb(0:n-1) = xb(0:n-1) - 2*x(0:n-1)*tempb
```

```
$ gfortran rosenbrock_b.f90 driverGrad.f90 -o gradient.exe  
$ ./gradient.exe  
Rosenbrock gradient: 15.9751 50.3524  
-109.0627 -447.2795 -260.4622 -113.7724  
-421.9504 -121.0793 66.6307 6.3868
```

Focus on forward mode: $F' \times \dot{V}_0$

$$F' \times \dot{V}_0 = f'_p(V_{p-1}) \times f'_{p-1}(V_{p-2}) \times \cdots \times f'_1(V_0) \times \dot{V}_0$$

\dot{V}_0 is a vector \Rightarrow compute from right to left!

This corresponds to P's original order

\Rightarrow interleave derivative and primal computation.

Easy!




Focus on reverse mode: $\overline{V_p} \times F'$

$$\overline{V_p} \times F' = \overline{V_p} \times f'_p(V_{p-1}) \times f'_{p-1}(V_{p-2}) \times \cdots \times f'_1(V_0)$$

Vector now on the left \Rightarrow compute from left to right

Not so easy, but worth the effort!




A **forward sweep**, and then a **backward sweep**

The derivative instructions form the **backward sweep**
“**Data-flow reversal**” to get V_k ’s in **reverse order**

Cost model

$$F : \text{in} \in \mathbb{R}^n \rightarrow \text{out} \in \mathbb{R}^m$$



- ▶ full F' cost grows like n using the forward mode
Good if $n \leq m$
- ▶ full F' cost grows like m using the reverse mode
Good if $n \gg m$ (e.g. $m = 1$ for a gradient)

Alternatives to AD

- ▶ Symbolic differentiation, Continuous adjoint...
Uses equations, not code ;
Duplicates discretization & coding work
- ▶ Finite Differences
Robust wrt coding style & non-differentiable code ;
inaccurate (2nd order contributions) ; only forward mode
- ▶ Complex-step
Astute use of Complex arithmetics ;
similar to AD ; only forward mode

Tool landscape

- ▶ Various AD models \Rightarrow various AD tools:
 - Store derivatives stuck to primals or in separate variables ;
 - Tape all forward computation or only chosen values ;
 - Preaccumulate partials ; Prefer recomputation to storage
- ▶ Various languages \Rightarrow Various tool interfaces:
 - AD by separate tool or embedded in application language ;
 - AD in the compiler ; Differentiated code visible or not ;
 - AD on restricted language or DSL

End of basics

Let's look in detail !